



Journal of Statistical Software

MMMMMM YYYY, Volume VV, Issue II.

<http://www.jstatsoft.org/>

RCIndex: Querying C/C++ Code in R with libclang

Duncan Temple Lang
University of California at Davis

Abstract

The ability to get information about C/C++ code routines and data structures can allow us to do many things in an interpreted language such as R. We use **libclang**, a flexible, embeddable library, to develop R functionality to obtain and use information about native code. We describe the **RCIndex** package which provides high-level functionality to access and utilize this information and also lower-level approaches to query and manipulate other aspects of native source code. We describe how to use the package and scenarios in which it is useful. This functionality is infrastructural and some of the functionality of the package is the native code analogy to the **codetools** package for R code. It is not necessarily of direct interest to end-users, but it allows us to build numerous tools that can help end-users and developers.

Keywords: R, **RCIndex** package, compiled code, meta-data, bindings.

1. Motivation

We typically think of C and C++ code as something we write, compile and call. It is rarely the input to anything but the compiler. However, such code is a source of potentially useful information that we can exploit in statistical and scientific computing. In order to leverage this information, we need to be able to access it in a structured manner and in a form that we can compute on and in the programming environment in which we want to use it. **libclang** (Carruth, Christopher, Gregor, Korobeynikov, Kremenek, McCall, Rosier, and Smith 2007; Gregor 2010) has emerged as a powerful, industrial-strength library that provides valuable functionality for working with C/C++ source code and which we can embed in R (and many other languages).

There are several applications of being able to programmatically understand C/C++ code in a high-level language such as R.

Registering R-callable routines: When we write C/C++ code to use in an R package, it is

useful to explicitly register the routines that we can call via any of the `.C()`/`.Call()`/`.External()` interfaces. When the routines are registered, R can help to identify potential errors in calling them. R can detect an incorrect number of arguments for a routine, or that the types are incorrect, e.g., an integer vector when a numeric vector is expected. Registration also allows us to resolve the symbols just once rather than each time we call a routine, and it also allows us to use different symbols/names to refer to the routines. It is convenient to be able to run R code to identify the R-callable routines and to generate the registration information programmatically. As we change these routines, we can programatically update the registration information with little effort and ensure the information is synchronized.

Generating bindings to native libraries: Numerous R packages provide interfaces to existing C/C++ libraries. This typically involves manually creating two pieces of code. The first is an R-callable wrapper routine corresponding to each routine of interest in the third-party library. The second is a corresponding R function that invokes the wrapper routine, having coerced the R arguments to the appropriate form. This is often quite straightforward, but both time-consuming and error-prone. This makes for unnecessarily lengthy write-debug-test cycles. Instead, we'd like to be able to programmatically read the information about the third-party routines and data types and then generate all of the code. We want to minimize human intervention. If we could generate these “bindings” programmatically, the R programmer can spend time creating higher-level functionality using these primitives.

Dynamic calls to native routines: The `rdyncall` (Adler 2012) and `Rffi` (Temple Lang 2011) R packages avoid having to explicitly create the wrapper routines and R functions to interface to existing C routines. Instead, they both provide a dynamic mechanism to call arbitrary native routines. However, both approaches require a description of the target routines. Again, we want to obtain this information programmatically and then we can easily generate these descriptions and remove humans from the process.

Understanding third-party libraries interactively: When we interface to third-party libraries, we typically read documentation to identify and understand the important routines and data structures. In some situations, it can be convenient to find this information interactively within an environment such as R. Rather than reading static document, we can query the code for information such as how often a particular data type is returned by a routine or passed as an argument? or what idioms does the library use? We can use R's graphics capabilities to visualize the code and which routines call which other routines.

In-line documentation as comments: Often third-party native libraries contain documentation in comments adjacent to the corresponding routines and data structures. Accordingly, it is convenient to be able to easily access this documentation and potentially reuse it as R documentation for functions that interface to the routines, as we did for the **RCUDA** package (<http://github.com/duncantl/RCUDA>).

Compiling R code: Recently, we have been developing R facilities for compiling R code to native instructions to by-pass the R interpreter. This allows us to rewrite and translate R code so that it is essentially native code and can call other existing native routines, for example, in the R engine or standard C libraries. This results in significant speedup. However, to make this work, we need to know the signature – parameter types and return type – of these native, external routines. Again, if we can find this information programmatically, we greatly simplify and improve the entire process of generating the code.

Determining memory management and mutability of inputs and outputs: When we

call existing native routines, we often need to know whether we are in charge of the memory for the inputs or whether the called routine is responsible. Often, programmers omit important information about whether a parameter is modified within a routine or if it is constant. This is important information as it allows us to differentiate between an array of values passed as a pointer, or a scalar whose contents are changed. We would like to be able to analyze the body of the routine to be able to determine if and how it modifies its parameters so that we can avoid making copies of data, if possible.

Software as Data: While we may not think of code as data, analyzing software is an important field of research and industry. Software defines several networks related to i) which routines call which other routines, ii) the hierarchical structure of C++ classes, iii) which files `#include` other files, and so on. We can explore how these networks change over time, i.e., modifications to the code. We can find which global/non-local variables are used in which routines to help identify problems with parallelization and potential refactorization of the code. We can combine this data with version control history to better understand software projects.

Detecting errors in native code for R: R's package mechanism provides a powerful set of tests and checks for potential errors in R code. These are very useful and can identify errors such as misspelled variable or function names before the code is run. It would be valuable to perform analogous tests on C code in R packages. We might identify common issues such as not protecting allocated R objects from garbage collection. We could do this by analyzing uses of *PROTECT()* and *UNPROTECT()* calls and ensuring there is an appropriate correspondence. These are R-specific checks and will not be done by other code-analyzing software, e.g., the compiler.

We can also find “dead” routines that are never called by other routines and so help to reduce the code.

We hope these applications motivate the utility of being able to navigate native code in R and indicate that there are many more. In this paper, we describe how to use the R interface to **libclang** to find this kind of information needed in these types applications. We start next by describing the fundamental concepts of **libclang** in section 2. We follow this in section 3 by introducing the high-level functionality provided by the **RCIndex** package, e.g., getting the routines, data structures, C++ classes in a source file. We then discuss the lower-level functionality in the package on which these higher-level functions are based in section 3.2. We explore three reasonably comprehensive examples or case-studies in section 4 and then compare this approach with others that I and others have pursued in the past. Finally, we outline some future plans for the package and its use.

2. Concepts of libclang

Before we discuss the R facilities for working with **libclang**, it is useful to understand the basic concepts of the library. These include translation units, abstract syntax trees and cursors, types and tokens. **libclang** exposes other concepts, but those are not the important ones for our purposes.

Source code is arranged in files. A project may be made up of more than one related files. For example, we may define related routines for one task in one file and other routines in a separate file. Header files are used for declarations and pre-processor definitions that are

shared by several files. When we parse a file, the parser essentially reads the source and substitutes the content of any header files referenced by an `#include` directive directly into the source. The entity resulting from parsing the document is called a *translation unit* (TU).

libclang is a parser and represents a translation unit as a parse tree where the nodes are called *cursors*. It first breaks the text into tokens during the lexical analysis step. This breaks code of the form `dnorm(n, 1, sd)` into the separate tokens “dnorm”, “(”, “n”, “1”, “sd”, “)”. The parsing stage then maps the tokens into language concepts represented by different kinds of cursors. In this example, the concept is a call expression. Within the call entity is the reference to the routine being called – `dnorm()` – and then the three arguments to the call. The parenthesis and comma tokens disappear in the parsing stage as we move from the individual tokens to higher-level semantic meaning of the tokens. However, a cursor still has access to its original tokens and we can obtain these for a given cursor. This turns out to be important as **libclang** doesn’t differentiate between different types of binary operators, for example. Instead, we find the operation (=, +, -, etc.) from the tokens and so the concepts of both cursor and token are important.

libclang represents the different parse elements via a cursor (of class *CXCursor* in R). A cursor has a kind that identifies its nature or what concept it represents. There are many different kinds for a cursor (enumerated in the variable *CXCursorKind* in the **RCIndex** package) and they include *FunctionDecl*, *VarDecl*, *ParmDecl*, *CallExpr*, *IntegerLiteral*, *ClassDecl*, *MacroDefinition*, *EnumDecl*, *StructDecl*. Hopefully the names indicate the purpose of the cursor.

Along with its kind, a cursor has child cursors. These are the components such as the reference to the routine and the parameters in a call, the fields and methods in a C++ class or the left and right hand side of an assignment expression. Each of these child cursors has a kind and also, possibly, its own sub-cursors. As a result, we have a tree, or hierarchical structure. Figure 1 illustrates this hierarchy for a simple routine shown in figure 2, along with the kind of cursor and its expression. The translation unit is the container for the top-most or root cursor of the tree. This root cursor has the abstract kind *TranslationUnit*. Its children are the top-level elements of the source code file, e.g., global variables, routines, class definitions, pre-processor terms. Each of these has its own child cursors and so on.

This concept of a cursor tree is important when we examine how to extract information using the lower-level facilities in sections 3.2 and 4. We will traverse the tree, either one node at a time or directly by querying child cursors and their children.

In addition to cursors, **libclang** manages type information. In the C and C++ languages, every expression has a type, be it a variable declaration, a call to a routine, a binary operator, and so on. **libclang** associates a type with each cursor. Importantly, it ensures that there is a single type object describing each unique data type in the translation unit, and across related translation units. This allows us to reason about and resolve types quite easily.

Like a cursor, a **libclang** type has a kind such as *Int*, *Float*, *Double*, *Typedef*, *Pointer*, *Record* (for a **struct**), *Enum*. A type also has a name, a size (number of bytes) and other characteristics we can query.

A translation unit contains all the information from the corresponding source file. A single TU is often sufficient for our purpose of exploration and analysis. However, if the code in one translation unit refers to routines or variables in an other source file, we have to merge it with another translation unit to resolve those references. **libclang** uses an *index* as the

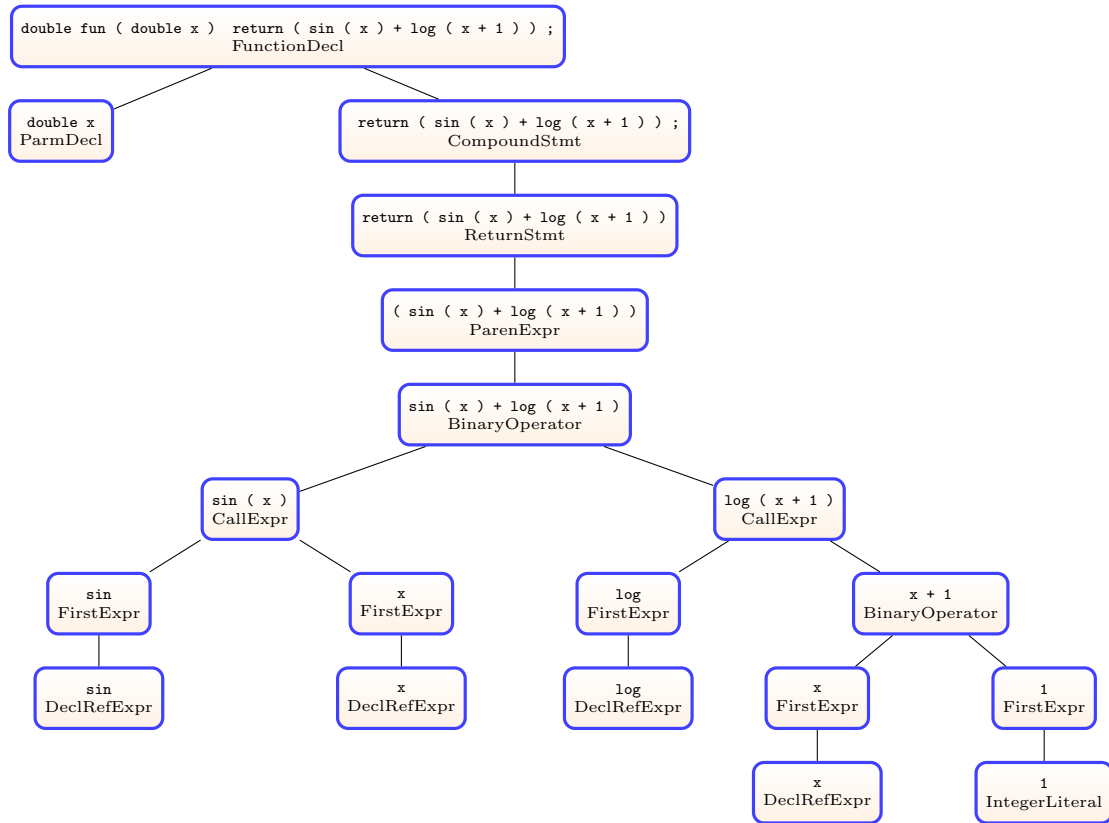


Figure 1: This shows the structure of a cursor tree for the C routine *sinLog()* shown in figure 2. Each node in the tree shows the kind of the cursor and the text of the expression with which it is associated. The cursor kind “FirstExpr” essentially means the cursor is considered opaque or hidden.

```

double
sinLog(double x)
{
    return(sin(x) + log(x + 1));
}

```

Figure 2: A simple routine that takes a single `double` value and returns a new value. Figure 1 illustrates the parse tree for this routine.

container for several related translation units. When we want to deal with the translation units together, we use the same index to parse them and this then allows **libclang** and us to connect references to elements across translation units.

libclang provides other features such as diagnostics for both parsing and compiler errors, code completion used in editors, serializing translation units, precompiled header files and efficient indexing of one or more translation units. Some of these are exposed in the **RCIndex** package. However, none are essential for understanding how to work with **libclang** and extracting useful information from source code.

3. The RCIndex package

The **RCIndex** provides numerous high-level functions that return information about C code. These refer to functions that I have developed building on the lower-level tools and which users can call to get useful information directly or to perform high-level tasks such as generating registration information or getting all the routines or C++ classes in a source file. These might be initial steps in a higher-level task, but these are high-level relative to the primitive functions in the **RCIndex** package with which we implemented these functions.

3.1. High-level Functionality

We might consider functions that take a source file and extract one or more of the types of top-level elements in that file, e.g., routines, data structures, enumeration definitions, C++ class definitions as the very highest-level functions. The R user can call these functions with just the name of the file and perhaps additional arguments for the parser and the results are returned. The user doesn't have to write any code to manipulate or traverse the parse tree (AST). She doesn't have to necessarily create a translation unit before calling one of these functions. As such, they are "atomic". These functions can also take an existing, previously parsed, translation unit rather than the file name. This is useful if we are going to make multiple separate queries of the same source code, i.e., we parse it once and query it multiple time.

The following paragraphs describe many of the high-level functions in the **RCIndex** package.

Finding routines The *getRoutines()* does as its name suggests. It takes a file name or an existing translation unit object and returns a list with an element for each routine declaration or definition in the corresponding code. Each element contains the *CXCursor* object representing the routine, a list of the parameters giving their name and data types, and the return type. This is an S4 object of class *FunctionDecl*. Since this contains the cursor, we can query it for the name of the routine, the file in which it is located, the location in that file and so on. In other words, information that we don't explicitly collect into the R object can be determined later when we use this description of the routine.

Data type definitions The routines are typically our focus so that we can invoke them from R or analyze their code. However, the routines work on data types and therefore we need to be able to find the definitions of these data types. The *getDataStructures()* function returns a list of the data types defined in a translation unit corresponding to a source file.

These include **struct** and **union** definitions, **typedefs** for providing an alternative name for an existing type and enumerated constant definitions.

C++ class definitions For C++ code, we can use *getCxxClasses()* to traverse a translation unit, either pre-parsed or a source file, and construct a description for each of the C++ classes it contains. Each class description is an instance of the class *CxxClass* in R, with template classes using the derived class *TemplateCxxClass*. A class description contains the name of the class, a list of its fields and another for its methods, and references to the base/super class cursors. The fields and methods each have their type information and also the accessor qualifier, i.e., **private**, **protected** or **public**. It is relatively straightforward to generate C++ code from this description that defines a new derived sub-class whose methods we can implement with R functions. We can then provide R instances of that class that operate transparently as regular C++ objects.

Enumerated constant definitions *getEnums()* returns a list of the enumeration definitions in a source file, i.e., corresponding to an expression of the form *enum { A, B, C }*. Each element is an instance of the S4 class *EnumerationDefinition*. Like the *FunctionDecl* class, this contains a reference to the definition in case we want to query it further at a later time. The actual values in the enumeration are stored in the **values** slot as a named integer vector. The names are the symbolic names we should use, while the values are the literal values to which these names correspond. These values allow us to cross the interface between R and native code where there is no existing association between the symbolic names and the literal values.

Global Variables The *getGlobalVariables()* function returns information about all of the top-level/non-local variables within a source file. From this, we have their names and types, and can also determine if they are constant, and if they are local to this file (i.e., **static**) or accessible to routines in other files. We may be interested in global variables for various reasons. We can use either of the **rdyncall** or **Rffi** packages to dynamically access the current value of a global variable. We also want to remove global variables when making code thread-safe or just for improving the logic of the code, i.e., avoiding side-effects.

The function *findGlobals()* finds global variables that are used within a routine or entire translation unit. The function is intentionally analogous to the *findGlobals()* function in **codetools** (Tierney 2011). It returns the variables, and optionally, the routines that are called, within a routine/translation unit which are not locally defined as parameters or local variables.

Find calls to routines The *findCalls()* function takes a cursor – typically either a routine or an entire translation unit – and finds the names of routines in any call made within this code. This allows us to determine which routines call which other routines and so describe a graph/network. We can then discover potentially interesting things about the routines. For example, we might look at just the routines within a single file and see which of these call other routines in the file and which are called by other routines in the file. This might identify isolated routines that don't necessarily belong in this particular file. It might also help us to understand how the routines interact and cluster. We can use the **igraph** (Csardi

and [Nepusz 2006](#)) and/or **graph** ([Gentleman, Whalen, Huber, and Falcon 2000](#)) packages to perform computations on the network and also to visualize it.

We'll show how to do this on a file **memory.c** in the R interpreter's source. We start by obtaining the routines from the file:

```
f = sprintf("%s/../../src/main/memory.c", R.home())
r = getRoutines(f, args = "-DHAVE_CONFIG_H",
               includes = c(sprintf("%s/../../src/include", R.home()),
                           sprintf("%s/include", R.home())))
```

We'll discuss the extra arguments **args** and **includes** in section [3.2](#).

For each of these routines, we find which routines they call with

```
kalls = lapply(r, findCalls)
```

To restrict the routines to only those within this file and create the adjacency matrix, we use

```
withinFileCalls = lapply(kalls, intersect, names(r))
m = matrix(0, length(r), length(r), dimnames = list(names(r), names(r)))
invisible(sapply(names(withinFileCalls),
                 function(id)
                   m[id, withinFileCalls[[id]]] <- 1))
```

We'll discard the routines that are not called by any other routine and themselves do not call any routine in the file. We do this with

```
i = rowSums(m) == 0 & colSums(m) == 0
m = m[!i, !i]
```

These are all basic R manipulations of R objects and have nothing to do with **libclang**.

Finally, we can draw the graph

```
library(igraph)
g = graph.adjacency(mm, "directed")
plot(g, vertex.size = 2, vertex.label.cex = .6, edge.arrow.size = .5)
```

and it is displayed in [figure 3](#).

I have used the **findCalls** function when creating bindings to third-party libraries to determine which routines currently have no wrapper routine.

Registering R-callable routines We mentioned in section [1](#) registering native routines we can call from R via the **.C()**, **.Call()** and **.External()** interfaces. The registration information is C code that looks something like

```
static R_CallMethodDef CallEntries[] = {
  {"R_csv_sample", (DL_FUNC) &R_csv_sample, 2},
  {NULL, NULL, 0}
```



```
};

void attribute_visible R_init_FastCSVSample(DllInfo *dll)
{
    R_registerRoutines(dll, NULL, CallEntries, NULL, NULL);
    R_useDynamicSymbols(dll, TRUE);
}
```

This example comes from the **FastCSVSample** package (<http://github.com/duncantl/FastCSVSample>). In this case, there is only one C routine we invoke via the `.Call()` function. In the **stats** package, there are 83 `.Call()` and 10 `.C()` routines, with registration information for all.

It is tedious, and hence error-prone, to manually create this registration information, and also to update it as the code evolves. Instead, we can use the `createRegistrationCode()` function to generate it for us. This is often as simple as something like

```
rg = createRegistrationCode("~/GitWorkingArea/FastCSVSample/src")
cat("#include <R_ext/Rdynload.h>", rg, sep = "\n", file = "init.c")
```

For an R package, we specify the `src/` directory of that package. We can also provide arguments for the parser such as **includes** or pre-processor definitions via the **args** parameter using `createRegistrationCode()`'s `...` parameter. The function processes each C and C++ file in the directory and determines which routines can be called via the `.C()` interface and which can be called by `.Call()`. It does this by examining the signatures and determining which are consistent with the two different call mechanisms. It then creates the registration information and also generates the `R_init_...()` routine, using the name of the package as suffix.

The R registration mechanism doesn't currently permit the programmer to specify which type of R object is expected for each parameter in a `.Call()` routine, e.g., an integer or numeric vector, or a function or a list. The **RCIndex** functionality does, however, allow us to infer this by analyzing the code within a routine and determining how it is being accessed. This could be added to the registration information and allow R to perform run-time checks or coercion to the appropriate types in call to the routine.

Listing included files The function `getInclusions()` allows us to obtain a list of all the files included in our translation unit, both directly and indirectly. This allows us to find what files include other files, which files are included by multiple files, what files are not included in a directory and so on. We can also visualize the implicit network.

3.2. Low-level building blocks

The functions we described above do not require the user to know any of the details of **libclang** and how **RCIndex** extracts the information from a source code file. To work with the results of some of these functions, however, the R programmer may need to know a little about **libclang**'s type system. Furthermore, to go beyond these functions and extract different information, a programmer needs to understand the lower-level tools on which these functions were built. At the very least, it is useful to know how to create a translation unit directly rather than re-parsing the same file multiple times. In this section, we discuss the functionality for creating a translation unit, working with cursor objects and cursor hierarchies. We'll address writing visitor functions after these topics.

Creating a translation unit One of the two vital functions in **RCIndex** is `createTU()`. This takes the name of a source file and parses it into a translation unit object in memory, returning a reference to that as an opaque R object. The signature of the function is

```
createTU(src, includes = character(),
        idx = createIndex(verbose = verbose),
        args = character(),
        verbose = getOption("ShowParserDiagnostics", TRUE), options = 0)
```

The **includes** parameter allows us to specify one or more paths to directories in which to find **#include**'d header files. We can pass arguments to control the parser via the **args** parameter. These can be pre-processor flags and definitions such as `-DHAVE_CONFIG_H`, or any of the options the **libclang** parser understands such as `-ferror-limit=1000` and `-fparse-all-comments`. The entire set of options is documented in the **clang** user manual (<http://clang.llvm.org/docs/UsersManual.html>)¹.

`createTU()`'s **options** parameter allows us specify certain additional control options for creating the parser beyond the **args** parameter. These control aspects such as skipping the bodies of routines when we just want the declarations, keeping a more complete record of the pre-processing steps, and other aspects that don't concern us. We specify these options either as a combination of bitwise-enumeration values from the `CXTranslationUnit_Flags` class. For instance, to skip the body of the routines and to indicate the translation unit is incomplete, we can use

```
CXTranslationUnit_SkipFunctionBodies | CXTranslationUnit_Incomplete
```

These two R variables represent these two options and we are combining them via the `|` operator. An alternative, but equivalent, approach is to use short-hand names as a character vector, e.g.,

```
c("SkipFunctionBodies", "Incomplete")
```

The **index** parameter for the `createTU()` function is often omitted as we typically parse a single source file and work on it separately from others. However, as we mentioned, routines often refer to other routines or variables in other source files in the same project. The index is a container for related translation units and provides the mechanism for resolving references across translation units it manages. Therefore, if we want to work on multiple related source files, we should first create an index and then pass this to each call to `createTU()`, e.g.,

```
index = createIndex()
tus = lapply(list.files("src", pattern = "\\.*c$", full.names = TRUE),
            createTU, index = index)
```

Alternatively, we can pass a vector of file names in the call to `createTU()` and it will create a translation unit for each, using the same index and **args** and **includes**. So our example could be written as

```
tus = createTU(list.files("src", pattern = "\\.*c$", full.names = TRUE))
```

¹The code generation options are not relevant as we are not generating code, only parsing it.

When creating an index, we can control whether it displays errors, warnings and general diagnostic information about the code on the console. We can disable this via the **verbose** parameter and passing `FALSE` as its value.

Working with cursors A translation unit is a container for the parse tree. We can query the name of the source file with `getFileName()`. More importantly, we can access the root cursor in a translation unit with either of the following expressions

```
as(tu, "CXCursor")
getTranslationUnitCursor(tu)
```

A cursor in R has the class *CXCursor*.

A cursor is a representation of a general semantic concept such as a call expression, a binary operation, a variable declaration and so on. We find out its particular kind or purpose using the `getCursorKind()` function or the short-hand `$` operator. So the following are equivalent

```
getCursorKind(cur)
cur$kind
```

Most cursors have an associated string such as the name of the routine being invoked in a call expression or a variable being referenced in a variable declaration. We get this string via the `getName()` function.

We can retrieve the actual text from the source file associated with a cursor using `getCursorTokens()`. This is important in some cases when the kind of the cursor doesn't give us enough detail to disambiguate between various possibilities. For example, a binary operator doesn't tell us what operator was being used. For this, we look at the actual text. `getCursorTokens()` conveniently returns the associated source code text already broken into lexical tokens with names that identify their token types, e.g., "Punctuation".

If we parse the code using the `-fparse-all-comments` argument to `createTU()`, we can retrieve the comments associated with certain kinds of cursors, i.e., variable, routine and type declarations. There are several functions to access the comment. We can get the text with `getRawCommentText()`. We can also get a brief version of the comment corresponding to certain conventions for marking up documentation in comments via `getBriefComment()`, i.e., the text to the end of the line after a `\brief` directive in the comment. The `getParsedComment()` function returns an actual *CXComment* object and we can treat this as a hierarchical object with child comments.

In some situations, it is important to map a cursor to its location in the source file. `getLocation()` returns the filename and the line and column numbers and the number of bytes (or offset) from the start of the file. **libclang** has various concepts of location and we can express which one we want to use by name, i.e., one of "Expansion", "Instantiation" and "Presumed". Some kinds of cursors are references to other elements in the translation unit. For example, in a call expression, the routine being called is represented by a reference cursor that refers to the actual routine. We can resolve that reference with the `getCursorReferenced()` function.

Recall that cursors are recursive structures, often with child cursors. We typically write functions that traverse this tree using `visitCursor()` and we will discuss this in the next section. However, it is convenient at times to explicitly access a cursor's children, and their

children and so on. We can conceptually think of the cursor as being like an R list with the children as its elements. This allows us to access the children individually using `cur[[i]]`, where i is the index of the child (with the first child at index 1). We can't use names to index the children or the `$` operator. We can find the number of children a cursor has using the `length()` function. We can access the entire list of children explicitly by calling the `children()` function and this allows us to loop over them with `lapply()` or `sapply()`.

In addition to being able to descend the cursor tree, we can navigate up the hierarchy using `getCursorSemanticParent()` and in a slightly different sense of the hierarchy with `getCursorLexicalParent()`.

Traversing a cursor tree To traverse a cursor tree, we typically use the `visitCursor()` function. This is (currently) more efficient than recursively processing the entire tree with R code looping over the children. However, writing visitor functions can be complicated as it typically involves maintaining state across calls to the same function. This is not very common in R so needs some discussion.

3.3. Writing visitor functions

The primary way to extract information from a translation unit is to use the `visitTU()` function. This takes the name of a source file or an already parsed translation unit. The second argument is a “visitor” function. This can take various forms but essentially corresponds to an R function (or the address of a native symbol). **libclang** iterates over some or all of the cursors in the translation unit and calls the R function for each each of these cursors, passing it both the current cursor and its parent cursor, for context. The function can perform any computations it desires and typically extracts and assigns any information it wants from each cursor, or the cursor itself.

The visitor function typically stores information that persists across the calls to it made by **libclang**. We then access these objects when **libclang** has completed traversing the tree. While we may be inclined to use global variables, e.g., in R's global environment or interactive “work space”, this is a bad idea. Instead, we want to use variables that are local to this particular visitor function. We can do this in two basic ways – reference classes or closures, also known as lexical scoping. In fact, these two are very similar.

We'll focus on a very simple example which merely iterates over the sub-cursors and finds the name and location of each routine in the translation unit. We are looking for cursors at the top-level of the translation unit which have a cursor kind `CXCursors_FunctionDecl`. We ignore any other cursors.

We will store the name of each routine and also its line number. We have to put these into a non-local variable within a call to our visitor function and then be able to retrieve them. This is a closure. Creating a closure is quite simple but requires understanding how R finds variables. One approach to creating a closure is to define a function that returns one or more functions. We often call this a generator function. We can define our generator function something like

```
genRoutineLocations =
function()
{
```

```

locations = integer()

visitor = function(cur, parent) {
  if(cur$kind == CXCursor_FunctionDecl)
    locations[[ getName(cur) ]] <- getLocation(cur)$location["line"]

  CXChildVisit_Continue
}

list(update = visitor, locations = function() locations)
}

```

There are several things to note. Firstly, we have defined the *genRoutineLocations()* function. That is not what we will use to traverse the tree. Instead, we call it to get a new pair of functions returned in a list. We will pass the **visitor** element of that list to *visitCursor()*. Secondly, note the use of `<-` when we assign the line number to the element of our *locations* vector. This is a non-local assignment and the updated value will persist across calls to this function. Thirdly, note that the *locations* vector is defined in the body of *genRoutineLocations*, not within our visitor function. The visitor function merely modifies it. Fourthly, our second function in the returned list (**locations**) is an anonymous function also defined during a call to *genRoutineLocations()* and it returns the current value of the *locations* variable when it is called. This allows us to retrieve the results after they were added in the different calls to our visitor function. Lexical scoping and closures are described in various texts about R and [Gentleman and Ihaka \(2000\)](#) is the definitive reference.

We can now use our generator function and obtain the line numbers of the different routines. We do so with

```

f = system.file("exampleCode", "fib.c", package = "RCIndex")
col = genRoutineLocations()
visitCursor(f, col$visitor)
col$locations()

```

Our call to *genRoutineLocations()* returns the list of two functions. We pass the **visitor** element to *visitCursor()*, along with the name of the source code file. Then we retrieve the updated contents with the second function **locations**.

As an aside, note that in this particular case, we don't need to recursively descend through the cursors as the routines will all be immediate children of the translation unit. Therefore, our visitor function returns *CXChildVisit_Continue* to move to the next sibling and not down across its children. This saves times processing irrelevant cursors.

Closures are very powerful but confuse some R programmers. Reference classes may be more familiar, especially to those familiar with classes in C++, Java or Python. The idea is that we define a class that has methods that share variables. Our methods will be our **visitor** and **locations** accessor functions in the generator function above, and the fields or variables will be the *locations* numeric vector. We can use the same code, but aggregate it in a different way to define a reference class as:

```

RoutineLocationVisitor =
setRefClass("RoutineLocationVisitor",

```

```

fields = list(locations = "numeric"),
methods = list(getLocations = function()
                locations,
                visitor =
                  function(cur, parent) {
                    if(cur$kind == CXCursor_FunctionDecl)
                      locations[[ getName(cur) ]] <-
                        getLocation(cur)$location["line"]

                    CXChildVisit_Continue
                  })

```

Here, we explicitly identify the shared variable and the methods via the **fields** and **methods** parameters. The only real difference between this and our generator function is that we have to call our accessor function to retrieve the *locations* value *getLocations()* rather than *locations()*. This is merely to avoid have a field and a method with the same name.

We can now use this reference class in almost exactly the same way we used our generator function:

```

col = RoutineLocationVisitor()
visitCursor(f, col$visitor)
col$getLocations()

```

The *RoutineLocationVisitor()* function returned by *setRefClass()* creates a new instance of this reference class. We pass the *visitor()* method to *visitCursor()* and then call the function *getLocations()* to obtain the results.

You can use any reference class you want and then pass the visitor or update method to *visitTU()*. However, it can be convenient to define your reference class to extend the *RefCursorVisitor* class in the **RIndex** package. This allows you to pass the entire reference object to *visitTU()* and to get the results back directly. This is merely syntactic sugar to simplify the programming. It changes the code above to create the reference class, call *visitCursor()* and retrieve the results to the more succinct

```
visitCursor(cur, RoutineLocationVisitor)
```

Essentially, extending the *RefCursorVisitor* class allows *visitCursor()* to identify the visitor and the result functions. The only change we need to make when defining our reference class is to add the **contains** argument in the call to *setRefClass()*, i.e.,

```

setRefClass("RoutineLocationVisitor",
            contains = "RefCursorVisitor",
            fields = list(locations = "numeric"),
            methods = list(getLocations = function() locations,
                           .....))

```

Similar to the *RefCursorVisitor* reference class, we can use an S4 class – *S4CursorVisitor* – to combine the visitor and result function and pass the two together to *visitCursor()*. We typically create our two functions as before via a generator function, and then combine the two functions into a formal object with, for example,

```
col = genCollector()
visitCursor( cur, new("S4CursorVisitor", update = col$update, result = col$vars))
```

Again, the purpose of this *S4CursorVisitor* class is merely to allow *visitCursor()* to identify the two functions – the visitor function and the function to access the results.

Copying cursors The *visitCursor()* function, by default, makes a copy of each of the C-level cursor objects in each call to our visitor function. This, combined with the garbage collection mechanism in the package, ensures that if the visitor function assigns a cursor to a non-local variable so it persists after that specific call, it will remain valid. If the visitor function does not need the cursors after each call, it can avoid this unnecessary copying. We do this by passing `FALSE` for the value of the **clone** parameter of the *visitCursor()* function. If the visitor function wants to store a cursor, it can explicitly copy the cursor itself by calling the *clone()* function and assigning the new cursor to a non-local variable. Generally, it is safest to use the defaults and incur the slight overhead of cloning.

Controlling how libclang traverses the tree By default, *visitCursor()* will arrange for the visitor function to be called for each cursor in the cursor hierarchy. Sometimes we don't need to see every cursor, but perhaps just the top-level cursors of the translation unit or the parameters in a routine, but not its body. A visitor function can indicate to **libclang** whether to recursively process the sub-cursors of the current cursor being visited, or to skip the entire sub-hierarchy and move to the next sibling of the current cursor. Alternatively, a visitor function may determine that it has seen enough and doesn't need to process any more cursors, i.e., gracefully terminate the traversal of the tree. To do this, the visitor function returns any of *CXChildVisit_Recurse*, *CXChildVisit_Continue* or *CXChildVisit_Break*. The visitor function can return a different value on each invocation and so dynamically determine where to visit next.

As an example of controlling the traversal, consider that we want to process all of the nodes in a particular routine named *bar()* but no other routines or elements of the translation unit. We could define our visitor function something like

```
genVisitor =
function()
{
  inBar = FALSE

  update =
    function(cur, parent) {
      if(cur$kind == CXCursor_FunctionDecl) {
        if(inBar) {
          cat("quitting having reached the routine", getName(cur), "\n")
          return(CXChildVisit_Break)
        }

        if(getName(cur) == "bar") {
          inBar <- TRUE
        }
      }
    }
}
```



```

        return(CXChildVisit_Recurse)
    } else
        return(CXChildVisit_Continue)
    } else {
        # processing the cursors within the bar routine.
        print(cur$kind)
        return(CXChildVisit_Recurse)
    }
}
}

```

We can then invoke it with

```

f = system.file("exampleCode", "mutateArg.c", package = "RCIndex")
visitTU(f, genVisitor())

```

We keep a non-local variable state variable to see if we have already encountered the routine *bar()*. If we have and we see another routine, we terminate the traversal. If we see a routine that isn't *bar*, we skip to the next routine with *CXChildVisit_Continue*. If the routine is named *bar()*, we tell **libclang** to process the sub-cursors.

4. Extended Examples and Applications

In this section, we explore some more advanced examples of how to use the **RCIndex** package to get different information from a translation unit. One of the best sources of such examples is the set of high-level functions in the package itself, e.g., *getRoutines()*, *getCppClassClasses()*. In the package and in the example in this section, we have used the generator function and lexical scoping approach to implement a set of collector functions that gather the different types of information we want.

In section 1, we provided numerous motivating applications of being able to process native code. We will present partial/heuristic approaches to some of these. One aim of these examples is to illustrate how to traverse the translation unit and sub-cursors and work on the structure of the information. Another aim is to illustrate how to work with **libclang**'s type system.

4.1. Checking for garbage collection errors in native code for R

In this example, we will write code that can examine a C routine written to be called via R's *.Call()* interface and try to identify if there are possible errors related to ensuring R objects are not garbage collected prematurely. The R API uses the macros *PROTECT()* and *UNPROTECT()* to mark an object as being in use and stop it from being garbage collected and then to unmark a number of protected objects. For example, the C code in figure 4 creates two new R objects and protects them both, performs some computations that populate the objects and then unprotects both with a call *UNPROTECT(2)*. In contrast, the code in figure 5 does not protect the R object it creates. It is quite possible that after allocating *ans*, R will release that object when it allocates *names* in the next expression. At that point, the

memory is corrupted and errors and crashes are likely. In other cases, we might protect the R objects we create, but fail to unprotect them, or at least some of them.

```
SEXP
R_foo_correct(SEXP r_x)
{
    SEXP ans, names;
    int n = Rf_length(r_x);
    double *x = REAL(x);
    PROTECT(ans = NEW_NUMERIC(n));
    PROTECT(names = NEW_CHARACTER(n));
    for(int i = 0; i < n; i++) {
        char *str;
        REAL(ans)[i] = foo(x[i], &str);
        SET_STRING_ELT(names, i,
                        mkChar(str));
    }
    SET_NAMES(ans, names);
    UNPROTECT(2);
    return(ans);
}
```

Figure 4: This is a simple and correct C routine that protects and unprotects the R objects it creates.

```
SEXP
R_foo_incorrect(SEXP r_x)
{
    SEXP ans, names;
    int n = Rf_length(r_x);
    double *x = REAL(x);
    ans = NEW_NUMERIC(n);
    names = NEW_CHARACTER(n);
    for(int i = 0; i < n; i++) {
        char *str;
        REAL(ans)[i] = foo(x[i], &str);
        SET_STRING_ELT(names, i,
                        mkChar(str));
    }
    SET_NAMES(ans, names);
    return(ans);
}
```

Figure 5: This version does not protect the R objects and, hence, also doesn't unprotect them.

We will develop an R function that will attempt to check for common problems related to garbage collection. Our function takes the **libclang** cursor for a routine and will then traverse the cursors throughout that tree. It will find calls to known routines that allocate R objects (e.g., *Rf_allocVector()*, *NEW_INTEGER()*, *NEW_NUMERIC()*) and also calls to *Rf_protect()* and *PROTECT()*, and *Rf_unprotect()* and *UNPROTECT()*. A simple test for correct code is to count the number of allocations, the number of calls to *Rf_protect()/PROTECT()* and attempt to determine the value passed to calls to *Rf_unprotect()/UNPROTECT()*. If these don't match, we can indicate a potential error. This is a simple heuristic approach, but it identifies many common cases.

Since we will traverse over the sub-cursors of the routine, we need a visitor function. We'll use a closure, but again, we could use a reference class. We need variables in which we can record the number of calls to each of categories of routines that allocate, protect and unprotect objects. We'll call these variables *numAllocs*, *numProtectCalls* and *numUnprotectCalls*, respectively. For the first two categories of routines, we just increment the current value in the corresponding R variable. For *Rf_unprotect()*, we'll collect the argument for each call to *Rf_unprotect()*. This might be a literal value, e.g., 2 in our example, or a variable or a general expression which makes our task more complex.

We have now identified the basic strategy for our visitor function. Next, we need to specify the details for the different types of cursors we encounter. A call to any of our routines in

our categories of interest will be identified by a cursor with a kind `CXCursor_CallExpr`. We can get the name of the routine being called using `getName()` applied to this cursor. (There are other ways also, but this is the simplest.) Based on this name, we update the relevant counter or ignore the call.

If the call in the cursor is to `Rf_unprotect()` or `UNPROTECT()`, we have to arrange to get the call's first and only argument. We can use either of two approaches for this. We can set a flag in our closure to indicate we are processing the sub-cursors of a `Rf_unprotect()` call and then subsequent calls to our visitor function can check this and interpret the sub-cursors appropriately. Alternatively, we can explicitly traverse the children of the current call cursor within our visitor function to determine the argument. We'll use the former approach as it illustrates setting and unsetting state across calls. We'll illustrate the second approach later in this example.

We define our generator function as

```
R_AllocRoutineNames = c("Rf_allocVector", "NEW_NUMERIC", "NEW_INTEGER",
                        "NEW_LOGICAL", "NEW_CHARACTER", "NEW_LIST")
```

```
genProtectAnalyzer = function()
{
  numAllocs = 0
  numProtectCalls = 0
  numUnprotectCalls = numeric()

  inUnprotect = FALSE
  allocCounterVarName = ""
  unProtectParent = NULL

  update = function(cur, parent) {
    if(inUnprotect && identical(unProtectParent, cur) ) {
      unProtectParent <- NULL
      inUnprotect <- FALSE
    }

    k = cur$kind      # get the kind of this cursor
    if(k == CXCursor_CallExpr) {
      fn = getName(cur) # name of the routine being called
      if(fn == "PROTECT" || fn == "Rf_protect")
        numProtectCalls <- numProtectCalls + 1
      else if(fn == "UNPROTECT" || fn == "Rf_unprotect") {
        numUnprotectCalls <- numUnprotectCalls
        unProtectParent <- parent
        inUnprotect <- TRUE
      } else if(fn %in% R_AllocRoutineNames)
        numAllocs <- numAllocs + 1
    } else if(inUnprotect) {
      if(k == CXCursor_IntegerLiteral) {
        val = getCursorTokens(cur)["Literal"]
```

```

        numUnprotectCalls <- c(numUnprotectCalls, as.integer(val))
    } else if(k == CXCursor_FirstExpr) {
        allocCounterVarName <- getName(cur)
    }
}

CXChildVisit_Recurse # visit this cursor's children
}

list(update = update,
      info = function() {
        list(numProtectCalls = numProtectCalls,
              numUnprotectCalls = numUnprotectCalls,
              inUnprotect = inUnprotect,
              numAllocs = numAllocs)})
}

```

For the present, ignore the first expression in our visitor function, the *if()* expression. We'll discuss this below. Our code examines the kind of the cursor and if it is *CXType_CallExpr* it does one thing and otherwise does another. For a call, we check which category of interest it is in, or not and update the relevant variables. For an unprotect call, we set the variable *inUnprotect* so that subsequent calls to this function will recognize this and extract the information from the argument to the unprotect call.

If the cursor is not a call expression, then we check to see if *inUnprotect* is TRUE. If it is, then we attempt to interpret the current cursor as (part of) the argument to the unprotect routine. Our goal is to get the value of that argument. Our simplified version checks to see if the current cursor is a *CXCursor_IntegerLiteral* type or a *CXCursor_FirstExpr*. The former indicates that we have a literal integer value. Unfortunately, there is no function in the **libclang** API for us to access the associated value directly. Instead, we use the *getCursorTokens()* function. This gives us the text around the cursor, broken into tokens or atomic elements that make sense at the C/C++-level. The elements in the resulting character vector are named with their token types and so we can extract the value named "Literal". If the cursor is so-called *CXCursor_FirstExpr*, we get its name as the value of a variable. In reality, we would process its sub-cursors.

The logic is quite simple – deal with a call expression or, if *inUnprotect* is TRUE, deal with other kinds of cursors. This raises the question as to when *inUnprotect* is ever set back to FALSE. This brings us back to the initial *if()* expression in the function:

```

if(inUnprotect && identical(unProtectParent, cur) ) {
    unProtectParent <- NULL
    inUnprotect <- FALSE
}

```

The purpose of this is to ensure that we don't continue to collect information from other cursors after we have processed any call to *UNPROTECT()*. If we didn't have this, the variable *inUnprotect* would remain TRUE even after we have process the entire call expression to the unprotect routine. As a result, if there are C expressions in the body of the routine after

the call to *UNPROTECT()* and they contain literal values or *CXCursor_FirstExpr* cursors, we will continue to accumulate information from these as if they related to the *UNPROTECT()* call. Accordingly, we need to determine when we have finished examining the sub-cursors of the *UNPROTECT()* call. To do this, we record the parent cursor of the *UNPROTECT()* call in the variable *unProtectParent*. Each time our visitor is called, we can compare the current cursor to this cursor and if they are the same, we are back to a sibling of the *UNPROTECT()* call.

Setting and unsetting a state variable across calls to a visitor function can be complex and often requires clear thinking. We could have used the other approach which was to detect the call to *UNPROTECT()* and immediately determine its arguments using *children()*. For example, we could have added the code

```
arg = children(cur)[[2]]
if(arg$kind == CXCursor_IntegerLiteral)
  as.integer(getCursorTokens(arg)["Literal"])
```

This makes the visitor function a great deal simpler and also slightly faster as we can avoid traversing the sub-cursors again. However, firstly, this doesn't take into account that the information may not be as readily accessible via the immediate children but may be in their descendants. We could write a separate function to visit these and call it from our visitor. Secondly, sometimes we have to use state, in particular when the information we need to extract is not in sub-cursors but in sibling cursors and their descendants. This is not uncommon.

With our visitor function defined, we are now ready to use on some routines. Before we can check a routine, we have to read it into R. We do this with *getRoutines()* via

```
f = system.file("exampleCode", "protectIncorrect.c", package = "RCIndex")
r = getRoutines(f, f, includes = sprintf("%s/include", R.home()))
```

Note that we specified the include directories so that the parser could find **Rinternals.h** and the other R header files.

We create our visitor function by calling the generator function *genProtectAnalyzer()* and then we pass the visitor element to *visitCursor()*, along with the routine we want to check:

```
v = genProtectAnalyzer()
visitCursor(r$R_foo_incorrect, v$update)
v$info()
```

The output is

```
$numProtectCalls
[1] 0

$numUnprotectCalls
numeric(0)

$numAllocs
[1] 2
```

We can see there are two allocations and no calls to protect or unprotect R objects. In contrast, when we run this on the correct version, each of the three counts has a value of 2. As a final remark about this example, we could make this more general and sophisticated. A reasonably common idiom is to use a variable to count the number of calls to *PROTECT()* we make, incrementing it each time. Then we pass this to *UNPROTECT()*, e.g.,

```
int n = 0;
PROTECT(a = Rf_allocVector(...)); n++;
PROTECT(b = Rf_allocVector(...)); n++;
PROTECT(c = Rf_allocVector(...)); n++;
PROTECT(d = Rf_allocVector(...)); n++;

UNPROTECT(n);
```

It is easy to omit incrementing the counter in one or more cases. We can try to trace this by identifying the variable in the call to *UNPROTECT()* and then finding out where it is incremented and try to find cases where it is not. In our code above, this should be the adjacent sibling of the call to *PROTECT()*. This is an example of requiring state across calls. We should note that if we find the *UNPROTECT()* call uses a counter variable, we can make a second pass over the cursor tree to find out where it is used.

4.2. Finding the signature for a foreign function interface call

In this example, we focus on working with types in **RCIndex**.

As I mentioned in section 1, there are two packages (**Rffi** and **rdyncall**) that allow us to dynamically invoke arbitrary native routines within an R session without having to write or compile any additional wrapper code. Both require a description of the signature of the routine to be called. That is, we need the return type and the number and types of the parameters. We don't want the user to have to know or specify the signature, so we want to be able to compute it for them. **rdyncall** uses GCC-XML to get the information about the routines and generate the corresponding signatures. Here we'll show how to do it with **RCIndex** and directly from the results of *getRoutines()* applied to the source code describing the routine(s) of interest.

Consider the simple routine

```
#include <math.h>
int
square_sin(int *val, int len, double *ans)
{
    int i;
    for(i = 0; i < len ; i++) {
        double tmp = sin(val[i]);
        ans[i] = tmp * tmp;
    }
    return(len);
}
```

This takes a vector of integer values and computes the square of the sin of each element and returns the values in the array of `double` values. It returns the number of elements it processed.

We'll look at how we can invoke this routine from R using **rdyncall** and then how we can generate the appropriate signature programmatically. Basically, before we can invoke the routine, we have to compile the C code into a shared library, load that into R and find the address of the symbol `square_sin()`. We do this with a shell command

```
R CMD SHLIB rffi.c
```

and the R code

```
dyn.load("rffi.so")
sym = getNativeSymbolInfo("square_sin")$address
```

Next, we'll create some sample inputs:

```
x = 1:10
ans = numeric(length(x))
```

To invoke the routine, we use the function `.rdyncall()` function in the **rdyncall** package. The call is

```
v = .rdyncall(sym, "*ii*d)i", x, length(x), ans)
```

`.rdyncall()` takes care of marshalling the inputs and outputs between R and the native code.

The strange (and unaesthetic) string `"*ii*d)i"` is the description of the signature of the routine. It indicates that the first argument (`*i`) is a pointer to one or more `int` values, the second (`i`) is a simple `int` value specifying the number of elements in the first argument, and the third (`*d`) is a pointer to one or more `double` values. The `)` character separates the inputs from the return type and the latter in this case is `i` and so an `int`. Of course, we knew this signature by looking at the code, although we didn't know the actual form it took for **rdyncall**.

We'll write a function that uses **libclang** and **RCIndex** to create the signature for an arbitrary routine. We'll assume our function is passed a *FunctionDecl* object returned by `getRoutines()`. Recall that this contains the parameters and the return type. We can easily generate the signature with something like

```
makeSig =
function(r)
{
  args = sapply(r@params, makeType)
  rtype = makeType(r@returnType)
  paste0( paste(args, collapse = ""), rtype)
}
```

The important part of this function is the call to `makeType()` for each type in the signature and that is the function we must implement next. The `makeType()` function is passed either a *CXCursor* or a *CXType*. We want to deal with the *CXType* so must convert a *CXCursor* first. (We could of course use S3 or S4 methods to do this also.)

```

if(is(ty, "CXCursor"))
  ty = getType(ty)

```

Once we have the *CXType*, we can query the kind of type it is with *getTypeKind()*. Given the kind, we first check to see if it corresponds to one of the primitive C types. We make a list of these by combining the different enumeration values for *CXType*, e.g., *CXType_Double*, *CXType_Int* and so on. We can do this with

```
BasicKinds = CXTypeKind[3:24]
```

since the primitive kinds are ordered contiguously in the enumeration.

If the type does correspond to a basic type, we map it to the string that *.rdyncall()* expects for this type. These are listed in a table in the help page for *.rdyncall()*. “i” corresponds to *CXType_Int*, “I” corresponds to *CXType_UInt* for an unsigned integer, and so on. We express this map via the named R vector

```

rdyncallMap = c(
  v = CXType_Void,
  B = CXType_Bool,
  C = CXType_Char_U,
  C = CXType_UChar,
  c = CXType_Char32,
  c = CXType_Char_S,
  S = CXType_UShort,
  I = CXType_UInt,
  J = CXType_ULong,
  L = CXType_ULongLong,
  s = CXType_Short,
  i = CXType_Int,
  l = CXType_Long,
  l = CXType_LongLong,
  f = CXType_Float,
  d = CXType_Double
)

```

The value in the vector is the **libclang** type kind. The corresponding name is the **rdyncall** type identifier. We can match the **libclang** kind to the values in the vector and then return the name of the corresponding element. We do this with the simple function

```

mapBasicKindToRdyncall =
function(kind)
{
  i = match(kind, rdyncallMap)
  if(is.na(i))
    stop("don't recognize that type")

  names(rdyncallMap)[i]
}

```


If the type kind of the type passed to *makeType()* is not a basic type, we have to examine it further. We can compare the kind to *CXType_Pointer* to determine if it is a pointer, or call *isPointerType()* with the actual type object. If the type is a pointer type, we can get the type of element it points to with *getPointeeType()*. Here we have to treat the C string type *char ** specially since **rdyncall** does. We check if the pointee type is of kind *CXType_Char_S* and if it is, we return the **rdyncall** string “Z”.

If the pointee type is not a *CXType_Char_S*, then we check to see if it is a basic type again. If so, we will combine the type string for this basic type with a *** prefix, e.g., “*i”. If the type is not a basic type, then it is a pointer to a generic C type and we return “p”, as **rdyncall** requires.

We also have to deal with types that are not simply basic or pointer types. For example, we may have a type declaration such as

```
typedef long VecSize;
```

This is simply an alias or a new name for an existing type. **rdyncall** doesn’t care about the name, but only the actual type. To deal with this in **RCIndex**, we can use the canonical type to get at the actual type of the alias and pass this to a call to *makeType()*.

The entire *makeType()* function is defined as

```
makeType =
function(ty)
{
  if(is(ty, "CXCursor"))
    ty = getType(ty)

  kind = getTypeKind(ty) # or ty$kind

  if(kind %in% BasicKinds) {

    mapBasicKindToRdyncall(kind)

  } else if(kind == CXType_Pointer) {

    pty = getPointeeType(ty)
    if(getTypeKind(pty) == c(CXType_Char_S))
      "Z"
    else if((pk <- getTypeKind(pty)) %in% BasicKinds)
      sprintf("%s", mapBasicKindToRdyncall(pk))
    else
      "p"

  } else if(kind == CXType_Typedef) {
    makeType(getCanonicalType(ty))
  }
}
```

With the function defined, we can use it on the routines in **rffi.c**:

```
r = getRoutines(system.file("exampleCode", "rffi.c", package = "RCIndex"))
sapply(r, makeSig)
```

```
square_sin    string1    string2    typedefn
"*ii*d)i"    "Z)v"      "Z)v"      "l)v"
```

These correspond to the declarations

```
int square_sin(int *val, int len, double *ans);
void string1(char *x);
void string2(const char *x);
void typedefn(VecSize s)
```

The final signature illustrates mapping our **typedef** example to the **rdyncall** string “l” corresponding to a **long**.

We could go further with generating descriptions for use with **rdyncall** or **Rffi** to process **struct** and **union** types and C++ classes and so on.

4.3. Determining unmodified inputs to a routine

We mentioned in section 1 that it can be useful to determine if a routine modifies the memory pointed to by its arguments. Ideally, programmers would qualify parameters with a **const** declaration to indicate that it is not modified. However, when this is not provided we do not know if it was omitted or if the routine does actually change a parameter. If we knew that it did not, we could avoid making a copy of the inputs before calling it. Similarly, we can differentiate between a routine that takes a pointer to a scalar **int** and returns a value in that memory, and a routine that takes a pointer to a collection of integers but does not modify them. In this example, we will build a simple-minded mechanism to try to determine which parameters are modified within a routine. We will identify the common idioms used to change values. We could also extend this to determine the associated length variable of an array of values passed as a pointer.

In this example, we show an alternative approach to traversing nodes in the tree. We use a visitor but within the visitor explicitly subset the children and their children with a single function call.

Consider the two routines in figures 6 and 7. Both take a pointer to a collection of **int** values. The first changes each element of the array. The second merely reads the elements of the array. From the declaration, we cannot tell the difference. We will write a visitor function that traverses the code of the routine and identifies expressions which modify any of the pointer parameters. For simplicity, we will pass the names of the parameters of interest, however, our function could compute the parameters as it traverses the entire tree for the routine. We'll start with a simple visitor function, defined in a generator function to create the closure:

```

int
foo(int *x, int len)
{
    for(int i = 0; i < len; i++)
        x[i] = 2*x[i];
    return(len);
}

```

Figure 6: This routine modifies the memory pointed to by its *x* argument.

```

int
bar(int *x, int len)
{
    int total = 0;
    for(int i = 0; i < len; i++)
        total = 2*x[i];
    return(total);
}

```

Figure 7: Like *foo()* on the left, this routine also expects a pointer to a *int*, but it does not modify the contents of the memory pointed to by the argument.

```

update =
function(cur, parent) {
    k = cur$kind
    if(k == CXCursor_BinaryOperator && "=" %in% (toks <- getCursorTokens(cur)) )
        processLHS(cur[[1]])

    CXChildVisit_Recurse
}

```

The basic strategy is that for each cursor, we determine if it is an assignment operator. There is no simple function to do this in the **libclang** API, but we can combine different pieces of information to get the answer. We check the kind of the cursor is *CXCursor_BinaryOperator* and then use *getCursorTokens()* to find the actual operator as a token. If it is a “=”, we have indentified our condition. Using cursor tokens is a little less precise than querying a specific cursor. The tokens may include an additional element before and/or after our cursor that may be misleading. Indeed, we might be wise to find the binary operator by examining the tokens of the left-hand side of the binary operator and extracting the final element of its tokens.

Once we have a binary operator, we are interested in its left-hand side, i.e., the first child. Rather than setting a state variable and continuing to traverse the sub-cursors, we will explicitly process this first child. We may have to traverse its children, but we will do that with a separate visitor function – *processLHS()*. This may make the computations a little simpler. So in our top-level visitor, we access the first child of the current cursor with *cur[[1]]* and then call the function that will determine what is being modified.

Our next task is to define the *processLHS()* function. We could define another visitor function to process these sub-cursor, but instead we will manually traverse this cursor tree for the left-hand side of our binary operator expression. We have various different situations to consider. This left-hand side can be arbitrary C code. Examples of the entire expression include

```

x[i] = 2*x[i];
total = 2*x[i];
*(x + i) = 2*x[i];

```

```

y[ x[i] ] = 2*x[i];
*addr(y, x, i) = 2*x[i];
y[ f(x, i, y) ] = 2*x[i];

```

These illustrate assigning to an element, a variable, a unary operator indexing x , nested subsetting, calling a routine to get an address, and subsetting with the index determined by a call to a routine. The last two of these call other routines and we would have to analyze their code to determine what is being modified.

We'll focus on the first four examples above. We define a function that can handle each of these different cases. `processLHS()` will be called with the cursor corresponding to the left-hand side of the assignment. We check to see if this is an array subscript (`[]`), a variable reference (`total`), a unary operator (`*`), another binary operator (`x+i`), a call expression (`f(x, i, y)` or `addr(y,x, i)`) and so on. The function is defined as

```

processLHS =
function(cur) {
  k = cur$kind

  id = if(k == CXCursor_ArraySubscriptExpr) { # x[i], y[x[i]], x[foo()]
    processLHS(cur[[2]])
    getName(cur[[1]])
  } else if(k == CXCursor_UnaryOperator){
    if(cur[[1]]$kind == CXCursor_ParenExpr)
      return(processLHS(cur[[1]][[1]]))
    else if(cur[[1]]$kind == CXCursor_DeclRefExpr)
      getName(cur[[1]])
    else if(cur[[1]]$kind == CXCursor_CallExpr)
      calls <- c(calls, getName(cur[[1]]))
  } else if(k == CXCursor_FirstExpr || k == CXCursor_DeclRefExpr) {
    # recursively from *( x + i) with x and again with i
    getName(cur)
  } else if(k == CXCursor_BinaryOperator) { # x + i arising in *(x + i)
    return(sapply(children(cur), processLHS))
  } else if(k == CXCursor_CallExpr)
    return(calls <- c(calls, getName(cur)))

  if(!is.null(id) && id %in% paramNames)
    modifiedVars <- c(modifiedVars, id)
}

```

The logic is very complex. We deal with each of the kinds of cursors we mentioned above. In several cases, we recursively call the `processLHS()` function to process sub-cursors. Each of the bodies in our `if()-else` expression either returns the name of a variable from the cursors it processes or simply returns. This allows us to conclude the function by comparing the variable name to those in which we are interested. If it is, the function updates `modifiedVars`. Note that this function is defined within our generator function so that it has access to the `modifiedVars` variable and also so that the `update()` function can refer to it.

This function does “visit” the child cursors, but not in **libclang** sense. Instead, it actively extracts and processes them without returning control to the **libclang** iterator. This makes the logic a little easier to understand.

This visitor function identifies the basic assignments to parameters and records the variables they modify. It also identifies cases it cannot understand due to calls to other routines within the left-hand side expression. It adds these to the *calls* variable and if it has any elements, we know we are in situation that is not known.

The function does not handle simple cases of aliasing such as

```
void fun(int *x)
{
    int *y = x;
    y[0] = 10;
}
```

Here, we have a local variable that corresponds to x and then we modify its contents. We might extend our function to identify each of these cases.

These examples have been more involved than many tasks we might want to do with **RIndex** but they show what is possible.

5. Other Research and Approaches

I have explored how we can get information about C/C++ code in R for many years and have used various different approaches. **libclang** appears to be the most robust, stable and flexible so far. **clang** and **libclang** are very active projects and important elements of the modern tool chain on different platforms. This means they will continue to improve, evolve and be maintained. The API (Application Programming Interface) is intentionally stable. As we will see later, **libclang** does the hard work in resolving types uniquely and simplifying the complexities of native code. **libclang** allows us to readily associate concepts in the parse tree with precise locations in the text of the source code. **libclang** doesn't allow us to modify the parse tree and rewrite native code. However, there is a quite different and more advanced API related to **libclang** that does and some of the concepts will carry over to that, should we need those facilities in R.

In the past, I have used the GNU compiler (*GCC*) [Stallman *et al.* \(2012\)](#) to export information about the source code using its *-fdump-translation-unit* command-line argument. This outputs many of the details of the source code in a reasonably low-level format. I developed the **RGCCTranslationUnit** [Temple Lang \(2009\)](#) package to both parse the output from *GCC* and also to collate the information into higher-level descriptions of routines, data structures, and generate binding code. Unfortunately, the format of *GCC*'s output is not well documented, and the information it exports has changed over releases. Generally, this is a valuable source of information, but not as stable or as readily usable or developed as **libclang**.

GCC-XML is an extension to *GCC* that, like the *-fdump-translation-unit* option, can export information about routines and data structures contained in C/C++ source code. As the name suggests, it exports in an *XML* format. We can then read this into R with, for example, the **XML** [Temple Lang \(1999\)](#) package. We can convert the *XML* content into descriptions of

the routines and types. Again, we have to implement this type registry. Also, GCC-XML does not give us access to the bodies of the routines and so we cannot address several of the applications we discussed in section 1. Since GCC-XML does not come with *GCC* itself, one has to install it separately. While this is not very difficult, it is another hurdle.

I also explored SWIG (Simplified Wrapper and Interface Generator) [Beazley \(2003\)](#) as a means to programmatically generate interfaces to native code. The idea was to leverage a widely used tool that could generate code for different languages. SWIG is written in PERL and was somewhat complex. It is more desirable to write the code generation mechanism in R itself rather than another language. This is because it is more familiar and also because more users can understand the code and contribute and extend the development. Also, as the name suggests, the focus of SWIG is to generate bindings to existing native code. It does not allow us to examine the body of the routines.

Many years ago, we adapted and embedded the little C compiler (lcc) [Hanson and Fraser \(1995\)](#) in the S language (both R and S-PLUS). This gave us many of the features we obtain with **libclang**, allowing us to access information about routines, data structures, etc. in source code. It does not support C++. Unfortunately, it did not readily deal with extensions to the C language used in the Linux header files. As a result, other than extending the code base, we couldn't continue to use it for our purposes.

Microsoft Windows provides facilities for accessing the information about routines and data structures in native code via its type libraries. I developed the **SWinTypeLibs** [Temple Lang \(2003\)](#) package to read this information generally. This gives us access to much of the information we need. Of course, it is specific to Microsoft Windows. It works on previously compiled code, not the source code. Accordingly, it gives us no access to the bodies of the native code.

All of the approaches are reasonable and have their advantages. However, **libclang** is a significant improvement over all of these other approaches for various different reasons. It provides most of the infrastructure we need and removes the burden for consumers to create this. It is flexible, but not excessively complex. It is tied to serious, ongoing projects and so continues to evolve. It allows us to work in the language in which we want to program, i.e., R. It provides access to all aspects of the native code, not just the declarations. It processes both C and C++.

The approach we described in this paper does not work when we don't have the source code available, but only compiled binary libraries. We will have header files available that provide the declarations of the routines and the data types of interest in the library. We can use these for many of the purposes. We cannot of course examine the bodies of the routines.

6. Future Work

We have created bindings for most of the facilities provided by **libclang**. However, we have ignored the code-completion routines and the various indexing facilities. Now that we have the basic tools provided by **RCIndex** in R, we can programmatically generate bindings to these routines.

Many of the bindings for **RCIndex** were generated programmatically. Accordingly, there is little documentation for these. We plan to adapt the documentation from the source code to provide more direct documentation in R.

The package provides a reasonably involved way to garbage collect the common C data types from **libclang** used in R, i.e., **CXCursor** and **CXTranslationUnit**. This is complicated as there is a dependency between these types as a **CXTranslationUnit** object cannot be released when it is no longer directly referenced in R. This is because there may be **CXCursor** objects it contains still referenced by R. We hope to extend this mechanism to **CXType** and **CXIndex**. Given these basic bindings, we plan to develop additional high-level functionality to process different aspects of native code. We are interested in processing the bodies of the routines to understand their characteristics. We also plan to extend the R code to generate bindings to routines and data structures. This code is currently in a sub-directory of the package (**generateCode/**) and so not directly exposed to R users. We also used this to generate many of the bindings for the **RCUDA** package (<http://github.com/duncantl/RCUDA>). We will merge the code generation mechanisms from the previous work on **RGCCTranslationUnit** and **SWIG**.

Over the last decade, some of us have considered making the R source code thread-safe and introduce user-level threads. While there are many reasons, both technical and non-technical, why this hasn't been done, it would be easier to do with tools that analyze the source code and identify where and how the code needs to change. While we don't have to develop these tools in R itself, it is convenient and **RCIndex** can be useful for developing such re-factoring tools.

While not exclusively related to **RCIndex**, we are experimenting with ways to compile aspects of the R language and R code to native machine code. We can then compile R visitor functions to make them significantly faster. Not coincidentally, the compilation uses LLVM as the code generator and LLVM is the back-end for **clang**, the C/C++ compiler from which **libclang** came.

While these are tasks that I think might be useful and I may work on, I welcome others to develop any of them separately, contribute and/or collaborate with me, or fork the package. The code is hosted at <https://github.com/omegahat/RClangSimple>.

References

- Adler D (2012). “rdyncall: Improved Foreign Function Interface (FFI) and Dynamic Bindings to C Libraries.” R package version 0.7.5, URL <http://CRAN.R-project.org/package=rdyncall>.
- Beazley DM (2003). “Automated scientific software scripting with SWIG.” *Future Generation Computer Systems - Tools for program development and analysis*, **19**(5), 599–609.
- Carruth C, Christopher E, Gregor D, Korobeynikov A, Kremenek T, McCall J, Rosier C, Smith R (2007). “libclang: C/C++ translation unit parser library.” http://clang.llvm.org/doxygen/group__CINDEX.html. <http://clang.llvm.org/docs/Tooling.html>.
- Csardi G, Nepusz T (2006). “The igraph software package for complex network research.” *InterJournal, Complex Systems*, 1695. URL <http://igraph.sf.net>.
- Gentleman R, Ihaka R (2000). “Lexical Scope and Statistical Computing.” *Journal of Computational and Graphical Statistics*, **9**, 491–508. URL <http://www.amstat.org/publications/jcgs/>.

- Gentleman R, Whalen E, Huber W, Falcon S (2000). *graph: graph: A package to handle graph data structures*. R package version 1.39.0.
- Gregor D (2010). “libclang: Thinking Beyond the Compiler.” llvm.org/devmtg/2010-11/Gregor-libclang.pdf.
- Hanson DR, Fraser CW (1995). *A Retargetable C Compiler: Design and Implementation*. Addison-Wesley.
- Stallman RM, *et al.* (2012). *Using the GNU Compiler Collection*. Free Software Foundation, Boston. <http://gcc.gnu.org/>.
- Temple Lang D (1999). “XML: An R package to parse and generate XML content.” <http://www.omegahat.org/RXML>.
- Temple Lang D (2003). “SWinTypeLibs: An R package to read Windows type libraries.” <http://www.omegahat.org/SWinTypeLibs>.
- Temple Lang D (2009). “Working with meta-data from C/C++ code in R: the **RGCCTranslationUnit** package.” *Computational Statistics*, **24**(2), 283–293.
- Temple Lang D (2011). “Rffi: Interface to libffi to dynamically invoke arbitrary compiled routines at run-time without compiled bindings.” R package version 0.3-0, URL <http://www.omegahat.org/Rffi>.
- Tierney L (2011). *codetools: Code Analysis Tools for R*. R package version 0.2-8, URL <http://CRAN.R-project.org/package=codetools>.

Affiliation:

Duncan Temple Lang
 4210 Math Sciences Building,
 University of California at Davis
 One Shields Avenue
 Davis, CA 95616
 E-mail: duncan@r-project.org
 URL: <http://www.omegahat.org>